

```
In [2]: import pandas as pd
df=pd.read_csv('HousingDB.csv')

In [3]: df

Out[3]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67	22.4
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88	11.9

506 rows × 14 columns

```
In [4]: df.isnull().sum()

Out[4]: CRIM      0
ZN          0
INDUS       0
CHAS        0
NOX         0
RM          0
AGE         0
DIS         0
RAD         0
TAX         0
PTRATIO     0
B           0
LSTAT       0
MEDV        0
dtype: int64

In [5]: df.columns

Out[5]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
              'PTRATIO', 'B', 'LSTAT', 'MEDV'],
              dtype='object')

In [6]: df.describe()

Out[6]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000

```
In [7]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM         506 non-null    float64
1    ZN           506 non-null    float64
2    INDUS        506 non-null    float64
3    CHAS         506 non-null    int64
4    NOX          506 non-null    float64
5    RM           506 non-null    float64
6    AGE          506 non-null    float64
7    DIS          506 non-null    float64
8    RAD          506 non-null    int64
9    TAX          506 non-null    int64
10   PTRATIO      506 non-null    float64
11   B            506 non-null    float64
12   LSTAT        506 non-null    float64
13   MEDV         506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB

```

```

In [8]: import seaborn as sns
import matplotlib.pyplot as plt
#To plot the graph embedded in the notebook
%matplotlib inline

```

```

In [9]: from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

```

```


In [26]: df['MEDV'].plot.box()

```

```

Out[26]: <Axes: >

```



```


In [30]: df=df[df['MEDV']<=35]
df['MEDV'].plot.box()

```

```


Out[30]: <Axes: >

```



```
In [31]: correlation_matrix=df.corr().round(2)
sns.heatmap(data=correlation_matrix)
```

Out[31]: <Axes: >




The correlation coefficient ranges from -1 to 1. If the value is close to 1, it means that there is a strong positive correlation between the two variables. When it is close to -1, the variables have a strong negative correlation

Look at the MEDV , it has strong correlation wiht RM

```
In [32]: plt.scatter(df['MEDV'],df['RM'])
plt.xlabel('House prices in 1000$')
plt.ylabel('RM')
```

```
Out[32]: Text(0, 0.5, 'RM')
```



From above we can see that, as House prices increase, RM also increases

```
In [33]: x=df[['RM']]
y=df['MEDV']
```

```
In [34]: print(x.shape)
print(y.shape)
```

```
(458, 1)
(458,)
```

Split data into 20:80 ratio. 80 will be training data and 20 will be testing data

```
In [35]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=5)
```

```
In [36]: model = LinearRegression()
model.fit(x_train, y_train)
```

```
Out[36]: ▼ LinearRegression
LinearRegression()
```

```
In [37]: y_pred=model.predict(x_test)
```

```
In [38]: mse=mean_squared_error(y_test,y_pred)
mse
```

```
Out[38]: 22.905315826290995
```

```
In [39]: import numpy as np
rmse=np.sqrt(mse)
rmse
```


```
Out[39]: 4.785949835329555
```

```
In [40]: r2_score=model.score(x_test,y_test)
r2_score
```

```
Out[40]: 0.2951091986432811
```

```
In [41]: plt.scatter(x_test,y_test,label='Actual Prices',color="green")
plt.plot(x_test,y_pred,label="Predicted Prices",color="black")
plt.title("Linear Regression")
plt.xlabel('RM')
plt.ylabel('Prices')
plt.legend(loc="best")
```

```
Out[41]: <matplotlib.legend.Legend at 0x29f24796250>
```



In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js